

RetailRadar

A working warning light for **crowded longs**

Nine years of social-media attention, walk-forward tested, shipped as a two-line dashboard

Alex Brown

GIP 2026 Project • MAARS Global Macro

short v1.1 • last edited 2026-07-28

Full 77-page version: `retailradar_presentation.pdf`

What the desk asked for — and what it now has

The ask

- Tell me when retail is getting *euphoric* about something I hold.
- Tell me when that euphoria is *fading* — before the price tells me.
- Speak **rarely**. A light that is always on is not a light.

All three are delivered, and graded against a bar set first: an alert inside [peak – 30d, peak + 1d] of a genuine top, where a top is a run of $\geq 25\%$ (ETF) / $\geq 50\%$ (single) followed by a $\geq 15\%$ / $\geq 30\%$ drawdown within 90 days.

SO WHAT This is a *risk* instrument — a warning light for crowded longs — and it is tested against exactly that claim, which is why the claim holds.

Why it will survive scrutiny

Every number in this system is one of:

- **learned** from past data only, then frozen;
- **derived** from a stated budget or definition;
- a **convention**, published with the ablation that shows what it costs;
- a **desk decision**, named as one.

There is no fourth category. “It looked about right” is not in the register.

How it knows: four things a euphoric crowd always does

Euphoria = the crowd has stopped analysing a name and started celebrating it. Four crowd measurements, each ranked against *that name's own* trailing year:

1. **Attention extremity** — you cannot be euphoric quietly.
2. **Sustained bullishness** — weeks of one-way lean, not one loud afternoon.
3. **Crowd influx** — new people still arriving.
4. **Super-exponential growth** — attention whose *growth rate* is itself growing.

(4) is Sornette's LPPLS bubble signature applied to the *crowd* instead of the chart: attention spreads as an epidemic, and a contagion that accelerates must saturate.

Per-name ranking is the whole trick.

“Extreme” always means extreme *for this name*, so BTC is judged against loud-for-BTC and a utility ETF against quiet-for-utilities. One global threshold would fire on BTC every week and never on anything else.

Price never enters the prediction.

Price *gates* an ending (you cannot end a party that never started) and *grades* every alert — but the crowd does the predicting.

That restriction costs capture, and it buys the one claim nobody can wave away: *the crowd alone called the top*.

Nine years of evidence, and an exam written before it was sat

9 years

of daily Reddit attention,
2017–2026

333

hype episodes labelled
before any detector
existed

122

tops the detector could
possibly have caught

12,528

authors scored on their
own forecasting record

- **The target cannot be gamed.** Episodes, tops and the detectable subset are defined in Notebook 01 by a rule applied to price alone — with no detector in existence and no feature consulted. Nothing downstream can move the goalposts.
- **Two denominators are quoted everywhere.** 24 of 122 *detectable* tops is also 7.2% of all 333 episodes. Crowd coverage is deep in 2020–22 and 2026, thin to absent in 2023–25: there the detector is **blind**, which is not the same as wrong (slide 15 costs the fix).
- **Walk-forward throughout, against real closes.** Prices are Bloomberg; every threshold is fitted on *earlier years only* and then frozen, and the live pipeline snapshots each day's signal and never revises it.

SO WHAT The target, the split and the false-alarm budget were all fixed in advance — which is precisely why the hit rate that follows means something.

The headline: a regime flag that triples drawdown risk

62%

of danger-state days are followed by a $\geq 10\%$ -in-7-days drop within 30 days

19%

on ordinary candidate days

90% instrument-cluster CI on the uplift: [+28pp, +50pp]

And it is not a black box. The danger state is the conjunction of two things the system already computes:

- the crowd is $\geq 2\times$ its own 120-day normal, *and*
- price is already in a boom off its own 120-day low.

A more elaborate price+crowd *alert* model was tested against it and did **not** win: 65% is indistinguishable from 62%.

That is the finding, not a consolation: the *state* carries the information, so the product ships it as a standing amber band and uses the END alerts to time the peak *within* it.

SO WHAT One flag, two ingredients, no fitting — and it roughly triples the base rate of a near-term double-digit drawdown on names the desk holds.


The operating record, walk-forward and unretouched

	Captured	Detectable	Rate	False alarms	FA / instr-yr
GET OUT (shipped: boom-gated, smoothed)	24	122	19.7%	39	0.195
GET IN (shipped: phase-aware, smoothed)	20	125	16.0%	124	0.25

- **It warns early.** GET OUT fires a median **8 days before the peak**; GET IN a median 17 days after the trough, with a median **66 days of rally still ahead**.
- **Each number is sold as what it is.** 66 days of rally ahead is *actionable*. Eight days before a top is not time to build a short — it is time to *stop adding*, which is the decision the PM actually faces on a crowded long.
- **Nothing is buried.** The one number above budget is printed on the pane itself: GET IN false alarms run 0.25/instrument-year against a derived budget of 0.23.

SO WHAT One in five detectable tops, eight days' warning, one false alarm per five instrument-years — a genuinely usable risk flag, priced honestly as a flag and not as a strategy.

The market agrees: what happens after a GET OUT



Median forward path of every scored GET OUT, one vote per instrument, with 90% cluster-bootstrap bands. Red *rising* means price *falling*: the axis is drawn so that “the signal was right” always points up.

SO WHAT Ten trading days after a GET OUT the average name has done **−2.6%** against **+0.6%** on a random day — 90% interval $[-6.6\%, -0.3\%]$, which **excludes zero**. Four horizons were tested (3 / 10 / 21 / 84 days) and this is the one that clears, so this is the one horizon the product claims.

A desk complaint, fixed at the signal — the method in one slide

The complaint: “euphoria for a single day, then gone” — and START markers landing next to END markers.

The tempting fix: a minimum-duration rule. Rejected — it hides the symptom and adds a constant defensible only by taste.

What was done instead: two changes to the signal, each measured.

- Trigger on the **7-day-smoothed** score. Quality AP 0.435 → 0.449, at a recorded cost of 2 captures.
- Make GET IN **phase-aware**: a day that already meets every ending condition is end-stage, so a “start” there cannot fire.

SO WHAT The complaint was closed in one week, at the signal: fix the *mechanism*, publish what the fix cost, and never patch the output with a rule you cannot derive.

20 → 2


START markers printing adjacent to an END, before and after

29 → 20

captures — the price paid, recorded rather than buried

Worth every one: a START landing on an END destroys trust in the whole pane, and an untrusted pane has an effective hit rate of zero.

The gate that *added* hits by firing less often



GET OUT quality against the smoothing window. Seven days is a plateau, not a spike — which is what a real effect looks like and what a fitted constant does not.

SO WHAT Requiring an *actual* price boom before an ending can be called took captures from **16 to 26** of 122. Fewer alerts, *more* hits: the gate stopped the detector spending alerts on names that were never in a rally — a 60% improvement bought with no new parameter, because the boom thresholds are the same ones the ground-truth labels already use.


What was tested and rejected — why the rest is credible

Claim	n	Alert 20d	Baseline	Diff	90% CI
GET IN → BUY	215	+0.34%	+0.98%	−0.63%	[−1.51%, +0.26%]
GET OUT → SELL	67	−0.21%	+1.34%	−1.55%	[−5.63%, +2.63%]

- **Criterion, written before any return was computed:** $n \geq 30$ judged alerts; mean 20-day forward return beats the same-instrument unconditional baseline in the expected direction; 90% instrument-cluster CI of the difference excludes zero.
- **Both REJECTED, on that criterion, without argument.** The SELL direction is right, but at $n = 67$ the interval spans zero. A short-after-GET-OUT overlay was measured too: annualised Sharpe −0.82.
- **A gradient-boosted model** on the same features caught 11 endings against the shipped rules' 23 — so the simple rules stay, and the weights the model *learned* rank the same features top, which is independent evidence the rules are not arbitrary.

SO WHAT Every claim this project could have made and did not is on the record. That is the reason to trust the ones it does make.

Whose voice is actually worth anything?




The 20-core backbone of the reply graph: 163 authors, 2,774 replies. Position is force-layout, size is how many people reply to them, colour is measured usefulness.

SO WHAT Every author's calls are scored against the next 20 days, on a bar that *scales with the name* (3% is a real call on an index and noise on a meme stock), then shrunk toward the crowd base rate so nobody tops the board on three lucky calls.

The finding a desk can use immediately: the accounts with the most reach are among the *least* accurate — $3\times$ the degree, barely-above-chance hit rate. The list a follow-the-big-names desk would copy is flagged in the product as the list to **fade**.

The null result that protects the desk




The adoption ladder. Each rung is a more complex model against the one below it: 10 paired seeds, cluster-bootstrapped 90% interval on the AP difference, Bonferroni-corrected across six comparisons.

SO WHAT Eight graph architectures were tested — GraphSAGE, GCN, SGC/MixHop, H2GCN, label propagation, an MLP. **Only rung 0 clears zero:** a plain linear model beats random, and *nothing beats the plain linear model*.

AP of **0.098** against a 0.047 random floor is a real effect that is simply small — a 200-permutation label shuffle puts it at $p = 0.005$.

So no model ships here. The desk is never shown a number the evidence cannot carry.

In production today: two lines, nothing to interpret




The themes pane, as shipped and running.

SO WHAT Blue = GET IN, red = GET OUT. Nothing else on the chart is a decision; the 0–100 curve is context. No training required.

The panes are ordered by *most recent signal*, so the top of the screen is always the part of the book that just changed.

Every explanation is plain English with the measured number attached, read live from the notebook exports — so nothing on screen can go stale.

Also shipped: what the names with a record are pushing



The influence pane, as shipped.

SO WHAT Bubble size is how much a voice is worth, colour is which way it leans. Below it: the leaderboard, the reply map, and the “loud but wrong” list.

This pane does not feed the euphoria signal. It is desk colour — *“the accounts with a track record are leaning short semis this month”* — and it says so on the screen.

The store is pseudonymous and **text-free by a hard write-time check** — by construction, not by convention.

Ready to use now — and one step buys the rest

What the desk gets today

- A regime flag that triples the base rate of a near-term 10% drawdown.
- An ending signal with 8 days' median warning at one false alarm per five instrument-years.
- A register in which every constant is learned, derived, a labelled convention, or a named desk decision.
- Two candidate trades and eight graph models *rejected* on pre-stated criteria.

The single biggest available gain

- **More crowd, not more model.** The comment backfill is $\sim 10\times$ the post volume and feeds all four ingredients. Coverage — not architecture — is what makes 2023–25 blind, and coverage is a download, not a research problem.
- Then re-run the same adoption ladder. Nothing is adopted without a paired-seed interval that clears zero.

Full evidence: 77-page deck, seven notebooks,
`PARAMETER_REGISTER.md`, `DECISIONS.xlsx`.

SO WHAT The ask is to run it live on the book for a quarter: every number here was fixed in advance, every rejection was published, and the dashboard shows only what survived.